

Predictive Modeling

Series 7

Exercise 7.1

In this problem, you will use support vector approaches in order to predict whether a given car gets high or low gas mileage based on the **Auto** data set.

- a) Load and inspect the data from **auto.csv**. When needed, define categorical features with dummy variables. Create a binary variable that takes on a 1 for cars with gas mileage above the median, and a 0 for cars with gas mileage below the median.
- b) Fit a support vector classifier to the data with various values of **cost**, in order to predict whether a car gets high or low gas mileage. Report the cross-validation errors associated with different values of this parameter. Comment on your results.
- c) Now repeat (b), this time using SVMs with radial and polynomial basis kernels, with different values of **gamma** and **degree** and **cost**. Comment on your results

Exercise 7.2

This problem involves the **OJ.csv** data set.

- a) Create a training set containing a random sample of 800 observations, and a test set containing the remaining observations.
- b) Fit a support vector classifier to the training data using **cost=0.01**, with **Purchase** as the response and the other variables as predictors. Inspect the different attributes of the classifier, in particular the number of support vectors, and the classes to which they correspond.
- c) What are the training and test error rates? Also give the corresponding confusion matrices.
- d) Use the **GridSearchCV()** function to select an optimal cost. Consider 10 values in the range 0.01 to 10. If the computation takes too much time, you may want to increase the tolerance by using the *tol* setting, or setting a maximum number of iterations using the *max_iter* setting.

- e) Compute the training and test error rates using this new value for **cost**.
- f) Repeat parts (b) through (e) using a support vector machine with a radial kernel. Use the default value for **gamma**.
- g) Repeat parts (b) through (e) using a support vector machine with a polynomial kernel. Set **degree=2**.
- h) Overall, which approach seems to give the best results on this data?

Result Checker

Predictive Modeling

Solutions to Series 7

Solution 7.1

a) **Python** code:

```
[1]: import numpy as np
import pandas as pd

# Load data
df = pd.read_csv('./data/auto.csv')

# Create new variable
df['mpg01'] = np.zeros(df.shape[0], dtype=int)
for i in range(df.shape[0]):
    if df.loc[i, 'mpg'] > df['mpg'].median():
        df.loc[i, 'mpg01'] = int(1)

# Redefine origin as a categorical variable
df = pd.get_dummies(data=df, drop_first=False,
                    columns=['origin'])
df = df.rename(columns={'origin_1': 'American',
                        'origin_2': 'European',
                        'origin_3': 'Japanese'})

# Drop and add columns
df = df.drop(['mpg', 'name'], axis=1)

# As a first inspection, print the first rows of the data:
print(df.head())
# As well as the dimensions of the set:
print('\nSize of Auto =\n', df.shape)
```

	cylinders	displacement	horsepower	weight	acceleration	year	mpg01	0
→ 8	8	307.0	130	3504	12.0	70	0	
1	8	350.0	165	3693	11.5	70	0	
2	8	318.0	150	3436	11.0	70	0	
3	8	304.0	150	3433	12.0	70	0	
4	8	302.0	140	3449	10.5	70	0	

	American	European	Japanese
0	1	0	0
1	1	0	0
2	1	0	0
3	1	0	0
4	1	0	0

Size of Auto =
(392, 10)

- b) The cross-validation errors can be produced with ease by the `GridsearchCV()` function from the `sklearn.model_selection` library. This is done by the code below.

```
[2]: from sklearn.model_selection import GridSearchCV
      from sklearn import svm

      n_folds = 10

      # Define predictors and response:
      y = df['mpg01']
      x = df.drop('mpg01', axis=1)

      # Set parameters to be tuned. Other options can be added
      costs = np.linspace(0.05, 100, 5)
      tune_parameters = {'C': costs}

      # Tune SVM
      clf_tune = GridSearchCV(svm.SVC(kernel='linear'),
                             tune_parameters,
                             cv=n_folds)

      clf_tune.fit(x, y)

      # Save Tune scores and corresponding standard deviation:
      error_tune = 1 - clf_tune.cv_results_['mean_test_score']
      error_std = clf_tune.cv_results_['std_test_score'] / np.sqrt(n_folds)

      # Print some attributes of model:
      print('Best parameter:\n', clf_tune.best_params_,
            '\nBest score:\n', np.round(1 - clf_tune.best_score_, 4))
```

```
Best parameter:
'C': 0.05
Best score:
0.0942
```

If we plot the values of the missclassification error, the minimum missclassification error occurs at cost 0.05.

```
[3]: import matplotlib.pyplot as plt

      # plot
      fig = plt.figure(figsize=(8, 5))
      ax = fig.add_subplot(1, 1, 1)

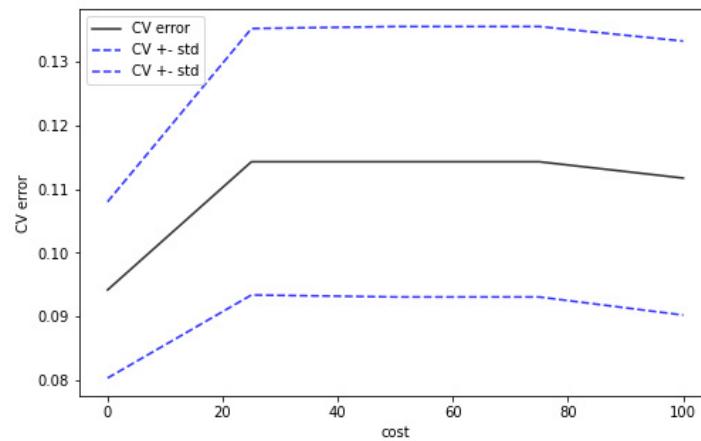
      ax.plot(costs, error_tune,
              '-k', alpha=0.8, label='Cross validation error')
      ax.plot(costs, error_tune + error_std, '--b',
              costs, error_tune - error_std, '--b',
```

```

alpha=0.8, label='Cross validation error standard deviation')

ax.set_xlabel('cost')
ax.set_ylabel('error')
plt.legend()
plt.show()

```



c) Radial Kernel:

```

[4]: # Set parameters to be tuned. Other options can be added
costs = np.linspace(0.5, 10, 5)
gamma = np.linspace(0.0005, 0.005, 5)

tune_parameters = {'C': costs,
                   'gamma': gamma}

# Tune SVM
clf_tune = GridSearchCV(svm.SVC(kernel='rbf'),
                       tune_parameters,
                       cv=n_folds)

clf_tune.fit(x, y)

# Save Tune scores:
error_tune = 1 - clf_tune.cv_results_['mean_test_score']
error_tune = error_tune.reshape(len(costs), len(gamma))

# Print some attributes of model:
print('Best parameter Radial:\n', clf_tune.best_params_,
      '\nBest score Radial:\n', np.round(1 - clf_tune.best_score_, 4))

```

```


Best parameter Radial:
  'C': 2.875, 'gamma': 0.0005
Best score Radial:
  0.0998

```

```
[5]: # plot
fig = plt.figure(figsize=(8, 5))
ax = fig.add_subplot(1, 1, 1)

# Plot error vs gamma for each value for cost:
for i in range(len(costs)):
    line, = ax.plot(gamma, error_tune[i, :],
                    '-.', alpha=0.8)
    line.set_label(('cost=' + str(costs[i])))

ax.set_xlabel('gamma')
ax.set_ylabel('CV error')
plt.legend()
plt.show()
```



Polynomial Kernel:

```
[6]: degree = [1, 2, 3, 4, 5]

tune_parameters = {'C': costs,
                  'degree': degree}

# Tune SVM
clf_tune = GridSearchCV(svm.SVC(kernel='poly'),
                       tune_parameters,
                       cv=n_folds)

clf_tune.fit(x, y)

# Save Tune scores:
error_tune = 1 - clf_tune.cv_results_['mean_test_score']
error_tune = error_tune.reshape(len(costs), len(degree))


# Print some attributes of model:
print('Best parameter Polynomial:\n', clf_tune.best_params_,
      '\nBest score Polynomial:\n', np.round(1 - clf_tune.best_score_, 4))
```

Best parameter Polynomial:
'C': 100.0, 'degree': 2
Best score Polynomial:
0.0996

```
[7]: # plot
fig = plt.figure(figsize=(8, 5))
ax = fig.add_subplot(1, 1, 1)

# Plot error vs degree for each value of cost:
for i in range(len(costs)):
    line, = ax.plot(degree, error_tune[i, :],
                    '.-', alpha=0.8)
    line.set_label(('cost=' + str(costs[i])))

ax.set_xlabel('degree')
ax.set_ylabel('CV error')
plt.legend()
plt.show()
```



d)

Solution 7.2a) **Python** code:

```
[1]: import numpy as np
import pandas as pd

# Load data
df = pd.read_csv('./data/OJ.csv')

# Redefine Store7 as a categorical variable
df = pd.get_dummies(data=df, drop_first=True,
                    columns=['Store7'])

# As a first inspection, print the first rows of the data:
print(df.head())
# As well as the dimensions of the set:
print('\nSize of OJ =\n', df.shape)
```

	Purchase	WeekofPurchase	StoreID	PriceCH	PriceMM	DiscCH	DiscMM	0	CH
→		237	1	1.75	1.99	0.00	0.0		
1	CH		239	1	1.75	1.99	0.00	0.3	
2	CH		245	1	1.86	2.09	0.17	0.0	
3	MM		227	1	1.69	1.69	0.00	0.0	
4	CH		228	7	1.69	1.69	0.00	0.0	

	SpecialCH	SpecialMM	LoyalCH	SalePriceMM	SalePriceCH	PriceDiff	0
→	0	0.500000	1.99	1.75	0.24		0
1	0	1	0.600000	1.69	1.75	-0.06	
2	0	0	0.680000	2.09	1.69	0.40	
3	0	0	0.400000	1.69	1.69	0.00	
4	0	0	0.956535	1.69	1.69	0.00	

	PctDiscMM	PctDiscCH	ListPriceDiff	STORE	Store7_Yes
0	0.000000	0.000000	0.24	1	0
1	0.150754	0.000000	0.24	1	0
2	0.000000	0.091398	0.23	1	0
3	0.000000	0.000000	0.00	1	0
4	0.000000	0.000000	0.00	0	1

Size of OJ =
(1070, 18)

```
[2]: # Set random seed
np.random.seed(1)

i = df.index
# Index of train
i_train = np.random.choice(i, replace=False,
                           size=800)

# Save DataFrames
df_train = df.iloc[i_train]
```

```
df_test = df.drop(i_train)

# Check dimensions:
print('\nSize of Train =\n', df_train.shape,
      '\nSize of Test =\n', df_test.shape)
```

```
Size of Train =
(800, 18)
Size of Test =
(270, 18)
```

b) **Python** code:

```
[3]: from sklearn import svm

# Define predictors and response:
y_train = df_train['Purchase']
x_train = df_train.drop('Purchase', axis=1)

# Fit SVM
cost = 0.01
clf = svm.SVC(kernel='linear', C=cost)
clf.fit(x_train, y_train)

# Print information on Support vectors:
print("Classes: ", clf.classes_,
      "\nNumber of Support Vectors: ", clf.n_support_)
```

```
Classes:  ['CH' 'MM']
Number of Support Vectors:  [312 313]
```

From the **Python** output we can see that 312 and 313 observations are used as support vector for the respective classes.

c) **Python** code:

```
[4]: # Define predictors and response:
y_test = df_test['Purchase']
x_test = df_test.drop('Purchase', axis=1)

# Find error rates:
e_train = 1 - clf.score(x_train, y_train)
e_test = 1 - clf.score(x_test, y_test)

print('Train error:\n', np.round(e_train, 4),
      '\nTest error:\n', np.round(e_test, 4))
```

```
Train error:
0.2175
Test error:
0.2222
```

```
[5]: def table_scores(ypredicted, ytrue):
    """ Return table showing predicted and true scores in n*n matrix
    Inputs:
    - Vector containing n predicted values
    - Vector containing n true values
    Returns:
    n*n Matrix with number of correct predictions on diagonal """
    # Empty Matrix:
    lables = np.unique(ytrue, return_inverse=False, axis=None)
    n = len(lables)
    scores = np.zeros((n, n))
    # Fill matrix with values:
    for i in range(len(ytrue)):
        true_class, pred_class = ytrue[i], ypredicted[i]
        scores[np.where(true_class == lables)[0][0]][
            np.where(pred_class == lables)[0][0]] += 1
    # Name rows and columns:
    r, c = [], []
    for i in range(len(lables)):
        r.append("True " + str(lables[i]))
        c.append("Pred " + str(lables[i]))
    scores = pd.DataFrame(scores, columns=c, index=r)

    return scores

print('Confusion Matrix train: \n',
      table_scores(clf.predict(x_train), y_train.to_numpy()),
      '\n\nConfusion Matrix test: \n',
      table_scores(clf.predict(x_test), y_test.to_numpy()))
```

Confusion Matrix train:

	Pred CH	Pred MM
True CH	432.0	49.0
True MM	125.0	194.0

Confusion Matrix test:

	Pred CH	Pred MM
True CH	151.0	21.0
True MM	39.0	59.0

d) **Python** code:

```
[6]: from sklearn.model_selection import GridSearchCV

n_folds = 10

# Set parameters to be tuned. Other options can be added
costs = np.linspace(0.01, 10, 10)

tune_parameters = {'C': costs}
```

```
# Tune SVM
clf_tune = GridSearchCV(
    svm.SVC(kernel='linear', max_iter=1e6, tol=1e-1),
    tune_parameters, cv=n_folds)
clf_tune.fit(x_train, y_train)

# Print some attributes of model:
print('Best parameter:\n', clf_tune.best_params_,
      '\nBest score:\n', np.round(1 - clf_tune.best_score_, 4))
```

Best parameter:

'C': 3.34

Best score:

0.1637

e) **Python** code:

```
[7]: print('Confusion Matrix train: \n',
        table_scores(clf_tune.predict(x_train), y_train.to_numpy()),
        '\n\nConfusion Matrix test: \n',
        table_scores(clf_tune.predict(x_test), y_test.to_numpy()))
```

Confusion Matrix train:

	Pred CH	Pred MM
True CH	427.0	54.0
True MM	71.0	248.0

Confusion Matrix test:

	Pred CH	Pred MM
True CH	148.0	24.0
True MM	24.0	74.0

f) **Python** code:

```
[8]: # Fit SVM with radial kernel
cost = 0.01
clf = svm.SVC(kernel='rbf', C=cost)
clf.fit(x_train, y_train)

# Print information on Support vectors:
print("Classes: ", clf.classes_,
      "\nNumber of Support Vectors: ", clf.n_support_)
```

Classes: ['CH' 'MM']

Number of Support Vectors: [319 319]

```
[9]: print('Confusion Matrix train: \n',
        table_scores(clf.predict(x_train), y_train.to_numpy()),
        '\n\nConfusion Matrix test: \n',
```

```
table_scores(clf.predict(x_test), y_test.to_numpy()))
```

Confusion Matrix train:

	Pred CH	Pred MM
True CH	481.0	0.0
True MM	319.0	0.0

Confusion Matrix test:

	Pred CH	Pred MM
True CH	172.0	0.0
True MM	98.0	0.0

```
[10]: # Tune SVM
clf_tune = GridSearchCV(svm.SVC(kernel='rbf'),
                        tune_parameters,
                        cv=n_folds)
clf_tune.fit(x_train, y_train)

# Print some attributes of model:
print('Best parameter:\n', clf_tune.best_params_,
      '\nBest score:\n', np.round(1 - clf_tune.best_score_, 4))
```

Best parameter:

'C': 0.01

Best score:

0.3988

```
[11]: print('Confusion Matrix train: \n',
          table_scores(clf_tune.predict(x_train), y_train.to_numpy()),
          '\n\nConfusion Matrix test: \n',
          table_scores(clf_tune.predict(x_test), y_test.to_numpy()))
```

Confusion Matrix train:

	Pred CH	Pred MM
True CH	481.0	0.0
True MM	319.0	0.0

Confusion Matrix test:

	Pred CH	Pred MM
True CH	172.0	0.0
True MM	98.0	0.0

g) **Python** code:

```
[12]: # Fit SVM with radial kernel
cost = 0.01
clf = svm.SVC(kernel='poly', C=cost, degree=2)
clf.fit(x_train, y_train)

# Print information on Support vectors:
```

```
print("Classes: ", clf.classes_,
      "\nNumber of Support Vectors: ", clf.n_support_)
```

```
Classes:  ['CH' 'MM']
Number of Support Vectors:  [319 319]
```

```
[13]: # Print confusion Matrix
print('Confusion Matrix train: \n',
      table_scores(clf.predict(x_train), y_train.to_numpy()),
      '\n\nConfusion Matrix test: \n',
      table_scores(clf.predict(x_test), y_test.to_numpy()))
```

```
Confusion Matrix train:
          Pred CH  Pred MM
True CH    481.0     0.0
True MM    319.0     0.0
```

```
Confusion Matrix test:
          Pred CH  Pred MM
True CH    172.0     0.0
True MM     98.0     0.0
```

```
[14]: # Tune SVM
clf_tune = GridSearchCV(
    svm.SVC(kernel='poly', degree=2),
    tune_parameters, cv=n_folds)
clf_tune.fit(x_train, y_train)

# Print some attributes of model:
print('Best parameter:\n', clf_tune.best_params_,
      '\nBest score:\n', np.round(1 - clf_tune.best_score_, 4))
```

```
Best parameter:
  'C': 0.01
Best score:
  0.3988
```

```
[15]: print('Confusion Matrix train: \n',
          table_scores(clf_tune.predict(x_train), y_train.to_numpy()),
          '\n\nConfusion Matrix test: \n',
          table_scores(clf_tune.predict(x_test), y_test.to_numpy()))
```

```
Confusion Matrix train:
          Pred CH  Pred MM
True CH    481.0     0.0
True MM    319.0     0.0
```

```
Confusion Matrix test:
          Pred CH  Pred MM
```

True CH	172.0	0.0
True MM	98.0	0.0

- h) The linear SVM performs best on this data. However, further investigating the optimal parameters for **gamma**, **degree** and **cost** could improve the behaviour of different classifiers.